

Report for Mean-Shift Algorithm

2.1 Implement the `distance` Function

We use 2-norm as the distance metric between points. Since radius is $+\infty$, we need to consider all the points. The code is as followed:

```
dist = torch.norm((X.float() - x.float()), dim=1)
```

2.2 Implement the `gaussian` Function

We compute weight as a Gaussian function of distance. The code is as followed:

```
weight = torch.exp(-(dist/bandwidth)**2/2)
```

Note that the normalization term of gaussian distribution would have no effect on the final output, since in the `update_point` or `update_point_batch` function, we will normalize the weights such that the sum of the weights is 1. So, it does not matter whether to normalize in `gaussian` function or not. To make the algorithm a bit faster, I don't use normalizer in the `gaussian` function.

2.3 Implement the `update_point` Function

To make algorithm faster, I don't use for loops. Instead, I transpose the weight and then multiply it by X matrix. Then the sum would be the new update point. The code is following:

```
weight = weight/weight.sum()
x = torch.sum(weight[:, None] *X, dim = 0)
```

2.4 Accelerating the Naïve Implementation

To accelerate the whole pipeline, we vectorize the inputs to avoid looping over each single point. The details are as followed:

- a) In `distance_batch` function, first reshape X for broadcasting, then we can compute the distances between any 2 points. Note that in this case, the first argument x is useless. The code is as followed:

```
d = X.reshape(-1, 1, 3) - X.reshape(1, -1, 3)
dist = torch.norm(d.float(), dim=2)
```

- b) In `update_point_batch` function, the weight is a matrix instead of a vector. Thus, I use matrix multiplication to compute the new points X_update:

```
weight = weight/(weight.sum(dim=1)[:,None]) # matrix
X_update = torch.matmul(weight.float(), X.float())
```

- c) In `meanshift_step_batch` function, call `distance_batch` function and `update_point_batch` function instead of `distance` function and `update_point` function.

Using both slow implementation (update for one point at a time) and fast implementation (update all points together using vectorization), we get following results:

- a) Slow implementation:

```
Elapsed time for mean-shift: 12.28938364982605
```

- b) Fast implementation:

```
Elapsed time for mean-shift: 2.279999256134033
```

As can be seen vectorization contributes to tremendous speed-up.